

Concurrent Programming

Mid-Sem Examination

1. Markup/programming hybrid

- Markup languages extended to support some programming
- Examples: XHTML, MXML (Action Script)

```
<mx:Button id="btn" label="MyButton" height="100" />
```

```
var btn:Button = new Button();  
btn.label = "MyButton";  
btn.height = 100;
```

2. Heterogeneous Programming:

1. For CPUs and accelerators like GPUs, FPGAs, Phis, ASICs
2. OpenCL, CUDA, OpenACC, OpenMP

3. Orthogonality

- A relatively small set of primitive constructs can be combined in a relatively small number of ways
- Every possible combination is legal
- C: function cannot return a static array, adding two pointers

Mid-Sem Examination

4. Tokens by Reg Exp, syntactic structure by CFG. Function parameters matching can be accepted by CFG. If-then-else cannot be accepted by CFG.
5. Shift-reduce parsers can accept left recursive grammars. Expression grammars as an example.
6. Binding ways - language design (operator), language implementation (float), compile time (variable), load time (static), runtime (dynamic)
7. Static scoping: $m.x=7, m.y=8, Q() \rightarrow Q.x=3, Q.y=4, P(4) \rightarrow m.x=6, \text{print}(Q.x/3), \text{print}(m.x/6) \rightarrow (3,6)$
Dynamic Scoping: $m.x=7, m.y=8, Q() \rightarrow Q.x=3, Q.y=4,$
 $P(4) \rightarrow Q.x=6, \text{print}(Q.x/6), \text{print}(m.x/7) \rightarrow (6,7)$
8. $E \rightarrow \text{id} \{ E.\text{type} = \text{id.type} \}$
 $E \rightarrow \text{int} \mid \text{float} \{ E.\text{type} = \text{int.type} \mid \text{float.type} \}$
 $E \rightarrow E1 + E2 \{ E.\text{type} = \text{resultant}(E1.\text{type}, E2.\text{type}) \}$

Mid-Sem Examination

```
E → id[ E1 ] { t1 = id.type;  
                if (t1 == ARRAY ^ E1.type == INTEGER)  
                    E.type = id.type;  
                else  
                    E.type = error;  
            }
```

Test Program

```
int x, y;  
function P(intn) {  
    x = (n + 2) * (n - 3)  
}
```

```
function Q() {  
    int x, y;  
    {  
        x = 3; y = 4;  
        P(y);  
        print (x);  
    }  
}
```

```
main() {  
    x = 7; y = 8;  
    Q();  
    print (x);  
}
```

Syllabus

Lecture Series (hours)	Topics
1-4	Introduction and Motivation, Paradigms
5-10	Syntax and Semantics, BNF, Compilation
11-18	Data Types, Constructs, Functions, Activation Records, Names and Bindings
19-28	Concurrency , Functional PLs, Logical PLs, Lambda Calculus, Event driven programming
29-36	Virtual Machines, Managed Languages, JIT, Case study

Concurrency

Two or more sequences of events occur “in parallel”

◆ Multiprogramming

- Single processor runs several programs at the same time
- Each program proceeds sequentially
- Actions of one program may occur between two steps of another

◆ Multiprocessors

- Two or more processors
- Programs on one processor communicate with programs on another
- Actions may happen simultaneously

Process: program running on a processor

The Promise of Concurrency

◆ Speed

- If a task takes time t on one processor, shouldn't it take time t/n on n processors?

◆ Availability

- If one process is busy, another may be ready to help

◆ Distribution

- Processors in different locations can collaborate to solve a problem or work together

◆ Humans do it so why can't computers?

- Vision, cognition appear to be highly parallel activities

Example: Rendering a Web page

- ◆ Page is a shared resource
- ◆ Multiple concurrent activities in the Web browser
 - Thread for each image load
 - Thread for text rendering
 - Thread for user input (e.g., “Stop” button)
- ◆ Cannot all write to page simultaneously!
 - Big challenge in concurrent programming: managing access to shared resources

The Challenges of Concurrency

- ◆ Concurrent programs are harder to get right
 - Folklore: need at least an order of magnitude in speedup for concurrent program to be worth the effort
- ◆ Some problems are inherently sequential
 - Theory – circuit evaluation is P-complete
 - Practice – many problems need coordination and communication among sub-problems
- ◆ Specific issues
 - **Communication** – send or receive information
 - **Synchronization** – wait for another process to act
 - **Atomicity** – do not stop in the middle and leave a mess

Language Support for Concurrency

◆ Threads (or process)

- Think of a thread as a system “object” containing the state of execution of a sequence of instructions
- Each thread needs a separate run-time stack

◆ Communication abstractions

- Synchronous communication
- Asynchronous buffers that preserve message order

◆ Concurrency control

- Locking and mutual exclusion
- Atomicity is more abstract, less commonly provided

Inter-Process Communication

- ◆ Processes may need to communicate
 - Process requires exclusive access to some resources
 - Process need to exchange data with another process
- ◆ Can communicate via:
 - Shared variables
 - Message passing
 - Parameters

Explicit vs. Implicit Concurrency

◆ Explicit concurrency

- Fork or create threads / processes explicitly
- Explicit communication between processes
 - Producer computes useful value
 - Consumer requests or waits for producer

◆ Implicit concurrency

- Rely on compiler to identify potential parallelism
- Instruction-level and loop-level parallelism can be inferred, but inferring subroutine-level parallelism has had less success

cobegin / coend

◆ Limited concurrency primitive

– Concurrent Pascal [Per Brinch Hansen, 1970s]

```
x := 0; y = 0;
```


```
cobegin
```

```
begin x := 1; x := x+1 end;
```

```
begin y := 2; y := y+1 end;
```

```
coend;
```

```
print(x);
```



Atomicity at level of assignment statement

Properties of cobegin/coend

- ◆ Simple way to create concurrent threads of execution (or processes)
- ◆ Communication by shared variables
- ◆ No mutual exclusion
- ◆ No atomicity
- ◆ Number of processes fixed by program structure
- ◆ Cannot abort processes
 - All must complete before parent process can go on

Race Conditions

- ◆ **Race condition** occurs when the value of a variable depends on the execution order of two or more concurrent processes

- ◆ **Example**

```
procedure signup(person)
begin
    number := number + 1;
    list[number] := person;
end;
signup(joe) || signup(bill)
```


Critical Section

- ◆ Two concurrent processes may access a shared resource
- ◆ Inconsistent behavior if processes are interleaved
- ◆ Allow only one process in critical section
- ◆ Issues
 - How to select which process is allowed to access the critical section?
 - What happens to the other process?

Locks and Waiting

<initialize concurrency control>

Process 1:

<wait>

signup(joe); // critical section

<signal>

Process 2:

<wait>

signup(bill); // critical section

<signal>

Need atomic operations to implement wait

Deadlock

- ◆ **Deadlock** occurs when a process is waiting for an event that will never happen
- ◆ Necessary conditions for a deadlock to exist:
 - Processes claim exclusive access to resources
 - Processes hold some resources while waiting for others
 - Resources may not be removed from waiting processes
 - There exists a circular chain of processes in which each process holds a resource needed by the next process in the chain
- ◆ Example: “dining philosophers”

Implementing Mutual Exclusion

◆ Atomic test-and-set

- Instruction atomically reads and writes some location
- Common hardware instruction
- Combine with busy-waiting loop to implement mutex

◆ Semaphore

- Keep queue of waiting processes
 - Avoid busy-waiting loop
- Scheduler has access to semaphore; process sleeps
- Disable interrupts during semaphore operations
 - OK since operations are short

Semaphores

- ◆ **Semaphore** is an integer variable and an associated process queue
- ◆ Operations:
 - $P(s)$ if $s > 0$ then $s--$
else enqueue process
 - $V(s)$ if a process is enqueued then dequeue it
else $s++$
- ◆ Binary semaphore
- ◆ Counting semaphore

Simple Producer-Consumer

```
program
  SimpleProducerConsumer;
var buffer : string;
    full : semaphore = 0;
    empty : semaphore = 1;
begin
  cobegin
    Producer; Consumer;
  coend;
end.
```

procedure **Producer**;

var tmp : string

begin

while (true) do begin

produce(tmp);

P(empty); { begin critical section
}

buffer := tmp;

V(full); { end critical section }

end;

end;

procedure **Consumer**;

var tmp : string

begin

while (true) do begin

P(full); { begin critical section }

tmp := buffer;

V(empty); { end critical section }

consume(tmp);

end;

end;

Producer-Consumer

```
program ProducerConsumer;  
const size = 5;  
var buffer : array[1..size] of string;  
    inn    : integer = 0;  
    out    : integer = 0;  
    lock   : semaphore = 1;  
    nonfull : semaphore = size;  
    nonempty : semaphore = 0; ...
```

```
procedure Produce  
var tmp : string  
begin  
    while (true) do begin  
        produce(tmp);  
        P(nonfull);  
        P(lock); { begin critical section  
    }  
        inn := inn mod size + 1;  
        buffer[inn] := tmp;  
        V(lock); { end critical section }  
        V(nonempty);  
    end;  
end;
```

```
procedure Consumer;  
var tmp : string  
begin  
    while (true) do begin  
        P(nonempty);  
        P(lock); { begin critical  
section }  
        out = out mod size + 1;  
        tmp := buffer[out];  
        V(lock); { end critical section }  
        V(nonfull);  
        consume(tmp);  
    end;  
end;
```

Monitors

- ◆ **Monitor** encapsulates a shared resource (monitor = “synchronized object”)
 - Private data
 - Set of access procedures (methods)
 - Locking is automatic
 - At most one process may execute a monitor procedure at a time (this process is “in” the monitor)
 - If one process is in the monitor, any other process that calls a monitor procedure will be delayed

Example of a Monitor

```
monitor Buffer;  
const size = 5;  
var buffer : array[1..size] of string;  
  in    : integer = 0;  
  out   : integer = 0;  
  count : integer = 0;  
  nonfull : condition;  
  nonempty : condition; ...
```

```
procedure put(s : string);  
begin  
  if (count = size) then  
    wait(nonfull);  
  in := in mod size + 1;  
  buffer[in] := s;  
  count := count + 1;  
  signal(nonempty);  
end;
```

```
function get : string;  
var tmp : string  
begin  
  if (count = 0) then wait(nonempty);  
  out = out mod size + 1;  
  tmp := buffer[out];  
  count := count - 1;  
  signal(nonfull);  
  get := tmp;  
end;
```

Java Threads

◆ Thread

- Set of instructions to be executed one at a time, in a specified order
- Special Thread class is part of the core language
 - In C/C++, threads are part of an “add-on” library

◆ Methods of class Thread

- start : method called to spawn a new thread
 - Causes JVM to call run() method on object
- suspend : freeze execution (requires context switch)
- interrupt : freeze and throw exception to thread
- stop : forcibly cause thread to halt

java.lang.Thread

```
public class Thread implements Runnable {  
    private char name[];  
    private Runnable target;  
    ...  
    public final static int MIN_PRIORITY = 1;  
    public final static int NORM_PRIORITY = 5;  
    public final static int MAX_PRIORITY = 10;  
  
    private void init(ThreadGroup g, Runnable target, String name) {...}  
  
    public Thread() { init(null, null, "Thread-" + nextThreadNum()); }  
    public Thread(Runnable target) {  
        init(null, target, "Thread-" + nextThreadNum());  
    }  
    public Thread(Runnable target, String name) { init(null, target, name); }  
  
    public synchronized native void start();  
  
    public void run() {  
        if (target != null) target.run();  
    }  
}
```

What does
this mean?

Creates execution environment
for the thread
(sets up a separate run-time stack, etc.)

Methods of Thread Class

```
public class Thread implements Runnable {
    ...
    public static native Thread currentThread();
    public static native void yield();
    public static native void sleep(long millis) throws InterruptedException;
    public static int enumerate(Thread tarray[])

    public static boolean interrupted() { ... }
    public boolean isInterrupted() { ... }
    public final native boolean isAlive();
    public String toString() {
    public void interrupt() { ... }
    public void interrupt() { ... }
    public final void stop() { ... }
    public final void suspend() { ... }
    public final void resume() { ... }
    public final void setPriority(int newPriority) {
    public final int getPriority() {
    public final void setName(String name) { ... }
    public final String getName() { return String.valueOf(name); }
    public native int countStackFrames();
    public final synchronized void join() throws InterruptedException {...}
    public void destroy() { throw new NoSuchMethodError(); }
}
```

Runnable Interface

- ◆ Thread class implements Runnable interface
- ◆ Single abstract (pure virtual) method `run()`

```
public interface Runnable {  
    public void run(); }
```
- ◆ Any implementation of Runnable must provide an implementation of the `run()` method

```
public class ConcurrentReader implements Runnable {  
    ...  
    public void run() { ...  
        ... code here executes concurrently with caller  
    ... }  
}
```

Two Ways to Start a Thread

- ◆ Construct a thread with a runnable object

```
ConcurrentReader readerThread = new ConcurrentReader();  
Thread t = new Thread(readerThread);  
t.start(); // calls ConcurrentReader.run() automatically
```

... OR ...

- ◆ Instantiate a subclass of Thread

```
class ConcurrWriter extends Thread { ...  
    public void run() { ... } }  
ConcurrWriter writerThread = new ConcurrWriter();  
writerThread.start(); // calls ConcurrWriter.run()
```

- ◆ What happens if you can just call
readerThread.run();

Why Two Ways?

- ◆ Java only has single inheritance
- ◆ Can inherit from some class, but also implement Runnable interface so that can run as a thread

```
class X extends Y implements Runnable { ...  
    public synchronized void doSomething() { ... }  
    public void run() { doSomething(); }  
}
```

```
X obj = new X();
```

```
obj.doSomething(); // runs sequentially in current thread
```

```
Thread t = new Thread(new X()); // new thread
```

```
t.start(); // calls run() which calls doSomething()
```

Interaction Between Threads

◆ Shared variables and method calls

- Two threads may assign/read the same variable
 - Programmer is responsible for avoiding race conditions by explicit synchronization!
- Two threads may call methods on the same object


◆ Synchronization primitives

- All objects have an internal **lock** (inherited from Object)
- **Synchronized method** locks the object
 - While it is active, no other thread can execute inside object
- Synchronization operations (inherited from Object)
 - Wait: pause current thread until another thread calls Notify
 - Notify: wake up waiting thread

Synchronized Methods

◆ Provide mutual exclusion

- If a thread calls a synchronized method, object is locked
- If another thread calls a synchronized method on the same object, this thread blocks until object is unlocked
 - Unsynchronized methods can still be called!



◆ “synchronized” is not part of method signature

- Subclass may replace a synchronized method with unsynchronized method

Wait, Notify, NotifyAll

```
public class Object {  
    ...  
    public final native void notify();  
    public final native void notifyAll();  
  
    public final native void wait(long timeout) throws InterruptedException;  
    public final void wait() throws InterruptedException { wait(0); }  
    public final void wait(long timeout, int nanos)  
        throws InterruptedException { ... }  
}
```

- ◆ **wait()** releases object lock, thread waits on internal queue
- ◆ **notify()** wakes the highest-priority thread closest to the front of the object's internal queue
- ◆ **notifyAll()** wakes up all waiting threads
 - Threads non-deterministically compete for access to object
 - May not be fair (low-priority threads may never get access)
- ◆ May only be called when object is locked (**when is that?**)

Using Synchronization

```
public synchronized void consume() {  
    while (!consumable()) {  
        wait(); } // release lock and wait for resource  
    ... // have exclusive access to resource, can consume  
}
```

```
public synchronized void produce() {  
    ... // do something that makes consumable() true  
    notifyAll(); // tell all waiting threads to try consuming  
    // can also call notify() and notify one thread at a time  
}
```

Example: Shared Queue

```
class SharedQueue {
    private Element head, tail;

    public boolean empty() { return head == tail; }

    public synchronized Element remove() {
        try { while (empty()) wait(); } // wait for an element in the queue
        catch (InterruptedException e) { return null; }
        Element p = head; head = head.next;
        if (head == null) tail = null;
        return p;
    }

    public synchronized void insert(Element p)
        if (tail == null) head = p;
        else tail.next = p;
        p.next = null;
        tail = p;
        notify(); // let one waiter know something is in the queue
    }
}
```

POSIX Threads

◆ Pthreads library for C

pthread_create - create a new thread giving it a “starting” procedure to run along with a single argument.

pthread_self - ask the currently running thread for its thread id.

pthread_join - join with a thread using its thread id (an integer value)

pthread_mutex_init - initialize a mutex structure

pthread_mutex_destroy - destroy a mutex structure

pthread_mutex_lock - lock an initialized mutex, if already locked suspend execution and wait

pthread_mutex_trylock - try to lock a mutex and if unsuccessful, do not suspend execution

pthread_mutex_unlock - unlock a mutex that was locked by the current thread

pthread_cond_init - initialize a condition variable structure

pthread_cond_destroy - destroy a condition variable structure

pthread_cond_wait - block the currently running thread on a condition variable indefinitely

pthread_cond_timedwait - block the currently running thread on a condition variable for a specific time

pthread_cond_signal - wakeup one thread blocked on a condition variable

pthread_cond_broadcast - wakeup all threads blocked on a condition variable

Example of Using POSIX Threads

```
#include <pthread.h>
#include <unistd.h> /* sleep declaration */
#include <stdio.h> /* printf declaration */
const int NUM_THREADS = 5;

void* sleeping(void* st)
{
    int sleep_time = (int) st; /* cast void* to an int */
    printf ("thread %d sleeping %d seconds ...\n", pthread_self(), sleep_time);
    sleep(sleep_time);
    printf ("\nthread %d awakening\n", pthread_self());
}

main( int argc, char *argv[] )
{
    pthread_t tid[NUM_THREADS]; /* array of thread IDs */
    int i;

    for ( i = 0; i < NUM_THREADS; i++)
        pthread_create (&tid[i], NULL, sleeping, i+2);

    for ( i = 0; i < NUM_THREADS; i++)
        pthread_join (tid[i], NULL);


    printf ("main() reporting that all %d threads have terminated\n", i);
} /* main */
```

Create several
child threads

Wait for children to finish

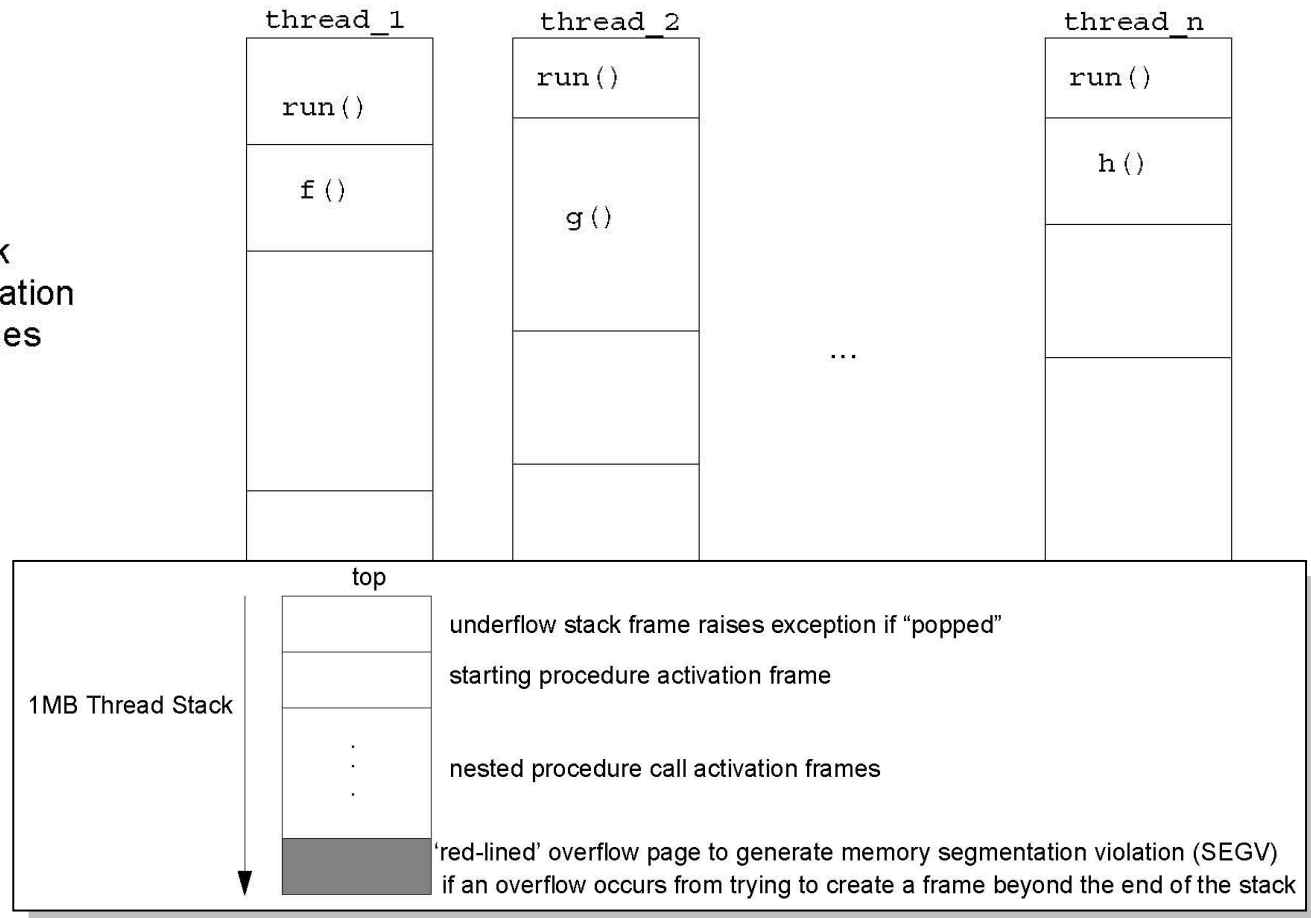
Thread Stacks

Main thread and run-time stack



Stack
Activation
Frames

Multiple thread run-time stacks, each a separate “thread of execution”



Thread Safety of Classes

- ◆ Fields of an object or class must always be in a valid state, even when used concurrently by multiple threads
 - What's a “valid state”? Serializability ...
- ◆ Classes are designed so that each method preserves state invariants on entry and exit
 - Example: priority queues represented as sorted lists
 - If invariant fails in the middle of a method call, concurrent execution of another method call will observe an inconsistent state

Making Classes Thread-Safe

- ◆ Synchronize critical sections
 - Make fields private, synchronize access to them
- ◆ Make objects immutable
 - State cannot be changed after object is created

```
public RGBColor invert() {  
    RGBColor retVal = new RGBColor(255 - r, 255 - g, 255 - b);  
    return retVal; }
```
 - Examples: Java String and primitive type wrappers Integer, Long, Float, etc.
 - Pure functions are always re-entrant!
- ◆ Use a thread-safe wrapper

Java-Style Synchronization in C++

```
class Synchronized {
    pthread_mutex_t m; // mutex variable
    pthread_cond_t c; // condition variable
protected:

    /* use this class to associate the mutex lock/unlock with the scope of a procedure */
    class Scope {
        Synchronized* obj;
    public:
        Scope(Synchronized* s) : obj(s) { pthread_mutex_lock(&obj->m); }
        ~Scope() { pthread_mutex_unlock(&obj->m); }
    };

public:

    Synchronized() { // initialize the mutex and condvar on construction
        pthread_mutex_init(&m, 0);
        pthread_cond_init(&c, 0);
    }

    ~Synchronized() { // destroy the mutex and condvar on destruction
        pthread_mutex_destroy(&m);
        pthread_cond_destroy(&c);
    }

    // map Java-like wait, notify and notifyAll onto pthread equivalents

    void wait() { pthread_cond_wait(&c, &m); }
    void notify() { pthread_cond_signal(&c); }
    void notifyAll() { pthread_cond_broadcast(&c); }
};
```

Thread-Safe Wrapper

- ◆ Define new class which has objects of original class as fields, provides methods to access them

```
public synchronized void setColor(int r, int g, int b) {  
    color.setColor(r, g, b);  
}  
public synchronized int[] getColor() {  
    return color.getColor();  
}  
public synchronized void invert() {  
    color.invert();  
}
```

Comparison

- ◆ Synchronizing critical sections
 - Good way to build thread-safe classes from scratch
 - Only way to allow wait() and notify()
- ◆ Using immutable objects
 - Good if objects are small, simple abstract data types
 - Benefits: pass without aliasing, unexpected side effects
- ◆ Using wrapper objects
 - Works with existing classes, gives users choice between thread-safe version and original (unsafe) one
 - Example: Java 1.2 collections library – classes not thread-safe, but some have methods to enclose objects in safe wrapper

Why Not Synchronize Everything?

◆ Performance costs

- Current Sun JVM – synchronized methods are 4 to 6 times slower than non-synchronized

◆ Risk of deadlock from too much locking

◆ Unnecessary blocking and unblocking of threads can reduce concurrency

◆ Alternative: immutable objects

- Issue: often short-lived, increase garbage collection

Syllabus

Lecture Series (hours)	Topics
1-4	Introduction and Motivation, Paradigms
5-10	Syntax and Semantics, BNF, Compilation
11-18	Data Types, Constructs, Functions, Activation Records, Names and Bindings
19-28	Concurrency, Functional PLs, Logical PLs, Lambda Calculus, Event driven programming
29-36	Virtual Machines, Managed Languages, JIT, Case study

Atomicity

[Flanagan]

An **easier-to-use** and **harder-to-implement** primitive:

```
void deposit(int x)
{
  synchronized(this)
  {
    int tmp =
    balance;
    tmp += x;
    balance = tmp;
  }
}
```

semantics:
lock acquire/release

```
void deposit(int x)
{
  atomic {
    int tmp =
    balance;
    tmp += x;
    balance = tmp;
  }
}
```

semantics:
(behave as if)

no interleaved execution

No fancy hardware, code restrictions, deadlock, or unfair scheduling (e.g., disabling interrupts)

AtomJava

[Grossman]

- ◆ New prototype from the University of Washington
 - Based on source-to-source translation for Java
- ◆ Atomicity via locking (object ownership)
 - Poll for contention and rollback
 - No support for parallel readers yet
- ◆ Key pieces of the implementation
 - All writes logged when an atomic block is executed
 - If thread is pre-empted in atomic, rollback the thread
 - Duplicate so non-atomic code is not slowed by logging
 - Smooth interaction with GC

Example: RGBColor Class

```
public class RGBColor {  
    private int r; private int g; private int b;  
    public RGBColor(int r, int g, int b) {  
        checkRGBVals(r, g, b);  
        this.r = r; this.g = g; this.b = b;  
    }
```

```
private static void checkRGBVals(int r, int g, int b) {  
    if (r < 0 || r > 255 || g < 0 || g > 255 ||  
        b < 0 || b > 255) {  
        throw new  
        IllegalArgumentException();  
    }  
}
```

What goes wrong with
multi-threaded use of this class?

```
public void setColor(int r, int g, int b) {  
    checkRGBVals(r, g, b);  
    this.r = r; this.g = g; this.b = b;  
}
```

```
public int[] getColor() {  
    // returns array of three ints: R,  
    int[] retVal = new int[3];  
    retVal[0] = r;  
    retVal[1] = g;  
    retVal[2] = b;  
    return retVal;  
}
```

```
public void invert() {  
    r = 255 - r; g = 255 - g; b = 255 -  
    b;
```

Problems with RGBColor Class

◆ Write/write conflicts

- If two threads try to write different colors, result may be a “mix” of R,G,B from two different colors

◆ Read/write conflicts

- If one thread reads while another writes, the color that is read may not match the color before or after